


DroneFollowingOrbsToPortal

Finetuning OrbFinder.yaml

- batch_size: 3000-> 1500
- batch_size: 1500-> 750
- batch_size: 750->375
- buffer_size: 60000->30000
- 4M -> agent found an imperfection in the environment and got stuck -> tried to get out -> started spinning in one place
- correction of imperfection -> at 5M it still doesn't find the portal with complete confidence
- between 5M and 7M the agent decides to clash into the wall -> rewards increased



The result with the final .yaml fájl



behaviors:

OrbFinder:

trainer_type: ppo

time_horizon: 256

max_steps: 2000000

hyperparameters:

learning_rate: 0.0003

batch_size: 750

buffer_size: 30000

num_epoch: 3

network_settings:

num_layers: 2

hidden_units: 256

reward_signals:

extrinsic:

gamma: 0.9

strength: 1.0

summary_freq: 2000